SOHAM SONAR

soham.sonar427@gmail.com ⬦ LinkedIn ⬦ United States ⬦ +1 (312) 975-7439 ⬦ GitHub ⬦ Portfolio

EDUCATION

Master of Computer Science — August 2023 - May 2025
Illinois Institute of Technology, Chicago, IL
Relevant Coursework: Machine learning, Cloud computing, Design and analysis of algorithms, Advanced operating systems

Bachelor of Computer Engineering — August 2018 - July 2022
Savitribai Phule Pune University
Relevant Coursework: Data structures and algorithm, Object oriented programming, Advanced database organization, Big data

EXPERIENCE

Research Assistant — February 2025 - October 2025
Gnosis Research Center - Illinois Institute of Technology — *Chicago, IL*

- Developed an integrated **Agentic AI platform** leveraging multi agent orchestration to automate end-to-end workflows across 40+ node clusters, enabling **autonomous task execution** and **intelligent operations coordination**.
- Enhanced the performance of **open source projects (IOWarp, Chronolog)**, by designing **REST APIs** and integrating an intuitive **assistant** for **data analytics** and **AI driven workflows**, reducing average data retrieval latency by 40%.
- Automated **CI/CD pipelines** with **GitHub Actions and Docker**, automating build, lint testing, and deployment processes across on prem systems and scalable cloud environments for faster and more reliable application delivery.
- Built **clean, reusable code** while researching LLM based applications **(Cursor, Claude)**, applying best practices in architecture, code reviews, unit testing, and scalability to ensure reliable **enterprise scale AI systems**.

Machine Learning Intern — January 2025 - April 2025
Vosyn Inc. — *Chicago, IL*

- Designed and optimized machine learning models using **Vertex AI, Kubeflow and Tensorflow** to improve real-time multilingual voice synthesis accuracy by 35%, ensuring seamless contextual translation across global markets.
- Integrated **10+ AI voice features** into customer facing applications through continuous model development and **A/B testing**, enabling real time support and improving usability for **non-technical users**.
- Deployed ML models for real-time voice localization using **Kubernetes & Cloud Run**, optimizing inference via **CPU/GPU benchmarking** reducing latency by 20% and enabled scalable cross-platform integration.

Executive — March 2023 - June 2023
Hexaware Technologies — *Mumbai, India*

- Architected and debugged software applications for healthcare solutions using **Python**, and **SQL**, achieving 60% performance improvement through **software architecture optimization**.
- **Streamlined data entry workflows**, reducing manual workload and improving data processing efficiency by 30% through scripting and **workflow automation**.
- Collaborated with **cross functional** teams in an **Agile Scrum** environment, and led backlog grooming and sprint planning across software engineering and QA teams, reducing post-deployment defects by 30%.

SKILLS

Programming Languages: Python, Java, C++, SQL, HTML, CSS, Bash/Shell Scripting.
AI Development: Autogens, FAISS, Finetuning, Github Copilot, Langchain, Vector DB, LLMs, RAG Models, Claude.
Machine Learning: Tensorflow, Pytorch, Scikit-learn, Huggingface, Keras, supervised & unsupervised models, Kubeflow.
Web/Software Development: Angular, React, Next.js, REST API, FastAPI, JIRA, Linux, Git, Agile, Scrum, VS Code.
Cloud/Big Data & Database: AWS, GCP, Docker, Kubernetes, Hadoop, Spark, Kafka, MySQL, MongoDB, PostgreSQL.

PROJECTS

Docquery [Link]

- Engineered an **AI powered knowledge** and document querying platform using **FastAPI, PostgreSQL, FAISS, and OpenAI embeddings**, enabling hybrid semantic–keyword retrieval with 30% higher answer accuracy than regular systems.
- Implemented scalable, production grade architecture with **Redis caching, async processing, supporting 200+ concurrent users** while reducing compute cost per query by 40% through intelligent caching and incremental indexing.

Enterprise IO Automation Framework [Link]

- Led the development of the Scientific **Model Context Protocol (MCP) server framework**, including Pandas, Parquet, Plot and HDF5 MCP servers, to **automate I/O** and **filesystem workflows** for local and cloud environments.
- Designed a **custom LLM client** using **Google Gen AI sdk** to coordinate **120+ simulation pipelines,** processing **multi-terabyte datasets** and significantly **reducing data access latency** across distributed systems.